

# AWS Assignment 1

## Objective:

This assignment will help everyone understand and implement key AWS services, including IAM, EC2, VPC, Subnets, and Nginx. The goal is to set up a basic web server in a secure VPC environment.

---

## Assignment Tasks:

### 1. IAM (Identity and Access Management)

- Create an IAM user with programmatic access and assign it to a custom IAM group with **EC2 and VPC Full Access**.
- Create a policy that allows the user to start, stop, and terminate EC2 instances but restricts access to other AWS services.

### 2. VPC (Virtual Private Cloud) and Subnets

- Create a **custom VPC** with CIDR block **192.168.0.0/16**.
- Inside the VPC, create two **subnets**:
  - **Public Subnet**: 192.168.1.0/24
  - **Private Subnet**: 192.168.2.0/24
- Set up an **Internet Gateway** and attach it to the VPC.
- Configure a **route table** to allow internet access only for the public subnet.

### 3. EC2 (Elastic Compute Cloud) Instance

- Launch an **EC2 instance** in the **public subnet** using Amazon Linux 2 or Ubuntu.
- Attach a **security group** that allows inbound SSH (port 22) and HTTP (port 80) access.

### 4. Install & Configure Nginx Web Server

- Connect to the EC2 instance via SSH.
- 

## Submission Requirements:

### 1. **Screenshots** of:

- IAM User & Policy
- VPC, Subnets, and Route Table Configuration
- Running EC2 Instance with Public IP

- Webpage running on Nginx
- 2. **Commands Used** (in a text file or PDF).
- 3. **Explanation** (brief write-up on what was learned).

Evaluation Criteria:

- ✓ Correct implementation of IAM, VPC, Subnet, EC2
- ✓ Successful Nginx installation and webpage hosting
- ✓ Clear documentation and screenshots

## AWS Assignment 2

Below is an example assignment that combines AWS Amplify, API Gateway, Lambda, SNS, and DynamoDB to build a simple serverless web application. This assignment is designed to help all gain hands-on experience with AWS serverless technologies.

### Objective

Build a serverless web application that allows users to submit data through a frontend application. The submitted data is processed by an API (via API Gateway and Lambda), stored in DynamoDB, and a notification is sent using SNS upon each successful data entry.

---

### Assignment Tasks

#### 1. AWS Amplify – Frontend Setup & Deployment

- **Create a New Amplify App:**
  - Initialize an Amplify project (using the Amplify CLI or Amplify Console).
  - Connect your Amplify project to a Git repository containing a simple static web app (e.g., built with React, Angular, or basic html css js).

- **Frontend Application:**
  - Develop a basic user interface with a form that collects sample user data (e.g., name, email, message).
  - Add functionality to call a REST API endpoint (to be created in Task 2) when the form is submitted.
- **Deployment:**
  - Deploy the frontend using Amplify Hosting.
  - Verify that the app is accessible via the Amplify-provided URL.

## 2. API Gateway – REST API Setup

- **Create a REST API:**
  - In the API Gateway console, create a new REST API.
  - Define a resource (e.g., `/submit`) with a POST method.
- **Integration with Lambda:**
  - Configure the POST method to trigger a Lambda function (created in Task 3).
  - Enable CORS on the API so that the Amplify-hosted frontend can call it.

## 3. AWS Lambda – Function Development

- **Create a Lambda Function:**
  - Develop a Lambda function in your preferred runtime (Node.js, Python, etc.).
  - The function should perform the following:
    - Parse the incoming JSON payload from API Gateway.
    - Insert the received data into a DynamoDB table (see Task 4).
    - Publish a notification to an SNS topic (see Task 5) confirming data receipt.
    - Return a suitable response (e.g., a success message and the stored data).
- **Permissions:**
  - Ensure the Lambda execution role has permissions to interact with DynamoDB and SNS.

## 4. DynamoDB – Data Storage

- **Create a DynamoDB Table:**
  - In the DynamoDB console, create a table (e.g., `UserSubmissions`) with an appropriate primary key (for example, `submissionId` as a UUID or timestamp).
- **Integrate with Lambda:**
  - Within your Lambda function, use the AWS SDK to insert new records into this table.
- **Data Validation:**
  - Test the table by manually inserting a sample record or using the Lambda function.

## 5. SNS – Notification Setup

- **Create an SNS Topic:**
  - In the SNS console, create a new topic (e.g., `SubmissionNotifications`).
  - Optionally, subscribe an email endpoint to the topic for real-time notifications.
- **Integrate with Lambda:**
  - Update your Lambda function to publish a message to the SNS topic after successfully inserting a record into DynamoDB.
  - The notification message should contain details about the new submission.

## 6. Testing & Validation

- **End-to-End Testing:**
    - Use the Amplify-hosted frontend to submit data.
    - Verify that:
      - The API Gateway correctly triggers the Lambda function.
      - The Lambda function stores the data in DynamoDB.
      - An SNS notification is published (and received if using email subscriptions).
      - A proper response is returned to the frontend, and the user sees confirmation.
  - **Debugging:**
    - Check CloudWatch logs for Lambda and API Gateway if issues arise during testing.
- 

## Submission Requirements

1. **Documentation:**
  - A report describing your architecture and each component's role.
  - A diagram illustrating the flow from Amplify to API Gateway, Lambda, SNS, and DynamoDB.
2. **Screenshots:**
  - Amplify Console showing the deployed frontend.
  - API Gateway configuration.
  - Lambda function code and CloudWatch logs.
  - DynamoDB table with sample data.
  - SNS topic configuration and any notification received (if applicable).
3. **Source Code:**
  - Provide a link to the Git repository containing your frontend application and Lambda function code.
4. **Commands & Configurations:**
  - A text file or PDF listing important CLI commands (if used) and configuration settings.

## Evaluation Criteria

- Correct integration of AWS Amplify, API Gateway, Lambda, SNS, and DynamoDB.
- Successful end-to-end functionality of the application.
- Clarity and completeness of documentation and code.
- Implementation of best practices for security (e.g., proper IAM roles and policies) and error handling.

# AWS Assignment 3 - Brief

Below is an assignment that integrates Step Functions, CloudWatch, SES, EventBridge Scheduler, EventBridge, SNS, and SQS into a cohesive serverless order processing pipeline. This scenario will give hands-on experience with orchestrating workflows, event-driven architectures, and monitoring/logging across multiple AWS services.

---

## Assignment: AWS Serverless Mini-Pipeline

### Objective

Build a simple serverless workflow that:

- Receives an order message via SQS.
  - Processes the message through a Step Functions workflow invoking a single Lambda function.
  - Uses that Lambda function to log the event in CloudWatch, send a confirmation email via SES, and publish a notification via SNS.
  - Uses an EventBridge Scheduler (or EventBridge rule) to trigger a heartbeat Lambda that logs a periodic message.
- 

### Assignment Tasks

#### 1. SQS – Simple Message Queue

- **Create an SQS Queue:**
  - **Name:** SimpleOrderQueue
  - Use default settings to keep it simple.
- **Send a Test Message:**

Use the AWS Console or CLI to send a sample JSON message (e.g.,  
`{ "orderId": "001", "customerEmail": "raju@gmail.com" }`  
```).

## 2. Lambda Function – Process Order Message

- **Create a Lambda Function:**
    - **Name:** ProcessOrderFunction
    - **Runtime:** (e.g., Python, Node.js)
  - **Function Tasks:**
    - **Log Event:** Write the incoming event details to CloudWatch.
    - **Send Email via SES:**
      - Use the SES API to send a simple confirmation email to the customerEmail received.
      - (Ensure that the sender email is verified in SES.)
    - **Publish SNS Notification:**
      - Publish a message to an SNS topic (created in Task 4) confirming order processing.
  - **IAM Permissions:**
    - Ensure the Lambda execution role includes permissions for CloudWatch logs, SES, and SNS.
- 

## 3. Step Functions – Minimal Workflow

- **Create a State Machine:**
    - **Name:** SimpleOrderWorkflow
    - **Definition:**
      - A single state that invokes the ProcessOrderFunction Lambda.
      - Pass the SQS message payload directly as input.
  - **Triggering the Workflow:**
    - **Option A:** Manually start an execution with the test message.
    - **Option B:** (For extra practice) Configure an EventBridge rule to trigger the state machine on a schedule.
- 

## 4. SNS – Notification Setup

- **Create an SNS Topic:**
    - **Name:** SimpleOrderNotifications
  - **Subscribe to the Topic:**
    - Add an email subscription (verify the email if necessary).
  - **Integration:**
    - The ProcessOrderFunction will publish a notification to this topic once the order is processed.
-

## 5. SES – Email Confirmation

- **Configure SES:**
    - **Verify Email:** Ensure your sender email is verified in SES (and the recipient if needed, due to sandbox restrictions).
  - **Integration:**
    - Within `ProcessOrderFunction`, use SES to send an order confirmation email containing basic order details.
- 

## 6. EventBridge Scheduler – Heartbeat Test

- **Create an EventBridge Rule or Scheduler:**
    - **Name:** `HeartbeatRule`
    - **Schedule:** Set a rule to trigger every 5–10 minutes (or a one-time test schedule).
  - **Create a Heartbeat Lambda Function:**
    - **Name:** `HeartbeatFunction`
    - **Task:** Log a simple “Heartbeat – Scheduler Triggered” message to CloudWatch.
  - **Integration:**
    - Configure the EventBridge rule to trigger `HeartbeatFunction` on schedule.
- 

## 7. CloudWatch – Monitoring & Logging

- **Enable Logging:**
    - Ensure all Lambda functions and the Step Functions state machine output logs to CloudWatch.
  - **Verification:**
    - Check CloudWatch logs to confirm that:
      - `ProcessOrderFunction` logs the received SQS message.
      - SES and SNS operations are logged.
      - `HeartbeatFunction` logs a scheduled heartbeat message.
- 

## Submission Requirements

1. **Documentation:**
  - A brief write-up explaining the purpose of each service and how they interact.
  - A simple diagram showing the flow:
    - `SQS` → `Step Functions` → `Lambda` → (`SES` & `SNS`)
    - EventBridge triggering the Heartbeat Lambda.
2. **Screenshots:**



- SQS Queue configuration with a sample message.
  - Step Functions state machine definition and a sample execution.
  - Lambda function code (or snippet) with CloudWatch log outputs.
  - SNS topic configuration and a screenshot of the received email notification.
  - EventBridge rule/scheduler configuration and the Heartbeat Lambda's CloudWatch logs.
3. **Source Code:**
- Provide code snippets for your Lambda functions.
4. **Deployment Instructions:**
- A brief guide on how to deploy and test your solution.
- 

## Evaluation Criteria

- **Service Integration:** Correct creation and integration of SQS, Step Functions, Lambda, SNS, SES, EventBridge, and CloudWatch.
  - **Functionality:** The pipeline should process a test message end-to-end, log events, send a confirmation email, and publish an SNS notification.
  - **Simplicity:** Implementation should be achievable within approximately 2 hours.
  - **Documentation:** Clear and concise documentation with diagrams and screenshots.
- 

This assignment allows you to gain hands-on experience with key AWS serverless and event-driven services, Happy building!

# AWS Assignment 4

Below is an assignment that combines Jenkins, Docker, GitHub, and shell scripting. This assignment is designed to help freshers understand how to create a CI/CD pipeline that builds and deploys a containerized application.

---

## Assignment: CI/CD Pipeline with Jenkins, Docker, GitHub, and Shell Scripting

### Objective

Create an automated CI/CD pipeline that:

- Hosts a simple web application in a GitHub repository.
  - Uses a Dockerfile to containerize the application.
  - Leverages Jenkins (running on an EC2 instance or your local machine) to trigger builds when changes are pushed to GitHub.
  - Uses Jenkins to build the Docker image, run a container, and execute shell scripts that verify the deployment.
- 

### Assignment Tasks

#### 1. GitHub Repository Setup

- **Create a Repository:**
  - Create a new GitHub repository (e.g., `simple-web-app`).
- **Develop a Simple Web Application:**
  - Create a minimal web application (e.g., using Node.js, Python Flask, or a static HTML page served by Nginx).
- **Dockerfile:**
  - Write a Dockerfile that:
    - Uses an appropriate base image.
    - Copies the web application code.
    - Installs necessary dependencies.
    - Exposes the appropriate port (e.g., 80 or 8080).
    - Specifies the command to run the application.

- **Push Code:**
    - Ensure the repository contains all the necessary code files, the Dockerfile, and a README with instructions.
- 

## 2. Jenkins Setup

- **Jenkins Installation:**
    - Install Jenkins on an EC2 instance or your local machine.
  - **Configure GitHub Integration:**
    - Set up a GitHub webhook so that Jenkins triggers a build when code is pushed.
  - **Create a Pipeline Job:**
    - Configure a Jenkins Pipeline (or freestyle job) that performs the following steps:
      - **Clone Repository:** Pull code from your GitHub repository.
      - **Build Docker Image:** Run `docker build` to create an image from your Dockerfile.
      - **Run Docker Container:** Use `docker run` to start a container from the built image.
      - **Execute Verification Script:** Run a shell script (see Task 3) to verify that the container is running and serving the web application.
- 

## 3. Docker & Shell Scripting

- **Docker Integration in Jenkins Pipeline:**
    - Ensure the Jenkins environment has Docker installed and the Jenkins user has permission to run Docker commands.
  - **Create a Shell Script (e.g., `verify.sh`):**
    - **Purpose:** Verify that the Docker container is up and the web application is responding.
    - **Tasks:**
      - Check if the Docker container is running (e.g., using `docker ps`).
      - Use `curl` to fetch the web page served by the container.
      - Print a success or error message based on the HTTP response.
  - **Integrate Shell Script:**
    - Include this shell script as a post-build step in your Jenkins Pipeline to confirm the deployment was successful.
- 

## 4. (Optional Bonus) Docker Hub Integration

- **Push Docker Image:**

- Extend your pipeline to tag the Docker image and push it to Docker Hub.
- Ensure that your Docker Hub credentials are securely stored in Jenkins (e.g., using credentials plugins).

## Submission Requirements

### 1. Documentation:

- A README file in your GitHub repository that explains:
  - The web application functionality.
  - How to build and run the Docker container locally.
  - Instructions on how the Jenkins pipeline is configured.
- A simple architecture diagram showing:
  - GitHub → Jenkins Pipeline → Docker (Build & Run) → Shell Script Verification.

### 2. Screenshots:

- GitHub repository structure.
- Jenkins pipeline configuration and build logs.
- Output of the shell script (e.g., successful curl response and container status).
- (Optional) Docker Hub repository view after pushing the image.

### 3. Source Code:

- Provide links to your GitHub repository containing:
  - Web application code.
  - Dockerfile.
  - Shell script (`verify.sh`).
  - Jenkins Pipeline script (Jenkinsfile) or configuration details.

### 4. Deployment Instructions:

- A step-by-step guide on how to set up the Jenkins job, configure the GitHub webhook, and run the pipeline.

---

## Evaluation Criteria

- **Integration:** Proper integration of GitHub, Jenkins, Docker, and shell scripting.
- **Functionality:** The pipeline should successfully clone the repository, build the Docker image, run the container, and execute the verification script.
- **Clarity & Documentation:** Clear instructions, comments, and diagrams that explain the process.
- **Automation:** Effective use of Jenkins to automate the build and deployment process.
- **Bonus Enhancements:** If Docker Hub integration is implemented, proper handling of image tagging and pushing.

This assignment will help you gain practical experience in setting up CI/CD pipelines with modern DevOps tools. Happy building!

# AWS Assignment 5

Below is a **small assignment** that combines Terraform, Linux, and Bash scripting. This task will have you provision infrastructure with Terraform, configure a Linux EC2 instance via a Bash user-data script, and then use an additional Bash script to verify and interact with the deployed instance.

---

## Assignment: Automated Web Server Deployment with Terraform & Bash Scripting

### Objective

Provision an AWS EC2 instance running Linux using Terraform. Automatically configure the instance to install and run a web server (Apache or Nginx) via a Bash user-data script, and then create a separate Bash script to SSH into the instance, verify system information, and check the web server's status.

---

### Assignment Tasks

#### 1. Terraform – Provision an EC2 Instance

- **Create Terraform Configuration Files:**
  - **main.tf:**
    - Configure the AWS provider.
    - Define an AWS EC2 instance resource with a suitable Linux AMI.
    - Create a security group that allows:
      - Inbound SSH (port 22)
      - Inbound HTTP (port 80)
  - **variables.tf:**
    - Define variables for instance type, region, key pair name, etc.
  - **outputs.tf:**
    - Output the public IP address of the EC2 instance.
- **Use a Bash User-Data Script:**
  - Embed a Bash script in the EC2 resource's `user_data` that:
    - Updates the system package index.
    - Installs a web server (choose either Apache or Nginx).

- Creates or modifies the default index.html to display a custom welcome message.
  - Starts and enables the web server service.
- 

## 2. Bash Scripting – Verification Script

- **Create a Bash Script (e.g., `verify.sh`):**
    - Use SSH (with the proper key) to connect to the newly provisioned EC2 instance.
    - Run commands on the instance to:
      - Display system information (e.g., `hostname`, `uptime`).
      - Check that the web server process is running.
      - Optionally, fetch the contents of the index.html (using `curl` or similar) to verify the welcome message.
  - **Make sure the script is executable** and document any prerequisites (such as SSH key configuration).
- 

## 3. Testing and Documentation

- **Deploy the Infrastructure:**
    - Initialize Terraform.
    - Run `terraform plan` to review changes.
    - Execute `terraform apply` to provision the EC2 instance.
  - **Run the Verification Script:**
    - Once the EC2 instance is up, run your `verify.sh` script to check the system information and web server status.
  - **Documentation & Diagram:**
    - Write a brief README that explains:
      - How to set up and run the Terraform configuration.
      - How the Bash user-data script configures the web server.
      - How to execute the verification Bash script.
    - Optionally, include a simple diagram showing:
      - Terraform provisioning → EC2 instance with user-data configuration → SSH verification.
- 

## Submission Requirements

1. **Terraform Code:**

- Include all Terraform configuration files (`main.tf`, `variables.tf`, `outputs.tf`, etc.).
  - 2. **Bash Scripts:**
    - Provide the user-data script (embedded in Terraform) and the separate `verify.sh` script.
  - 3. **Documentation:**
    - A README file with deployment instructions and an explanation of your setup.
    - A diagram (can be a simple image or ASCII diagram) showing the architecture flow.
  - 4. **Screenshots/Outputs:**
    - Screenshot of a successful `terraform apply` output.
    - Terminal output from running `verify.sh` showing system details and web server status.
- 

## Evaluation Criteria

- **Correct Terraform Configuration:**
    - Provisioning of an EC2 instance with the required security group and outputs.
  - **Effective Bash Scripting:**
    - Proper installation and configuration of the web server via the user-data script.
    - A working verification script that successfully connects via SSH and checks the instance.
  - **Documentation & Clarity:**
    - Clear instructions and a logical explanation of each component.
    - Inclusion of a diagram to illustrate the workflow.
  - **End-to-End Functionality:**
    - The deployed instance should be accessible (via its public IP) with the custom web page, and the verification script should confirm the server's status.
- 

This assignment will help you gain practical experience with infrastructure as code using Terraform, automate instance configuration with Bash scripting, and interact with Linux systems over SSH. Happy building!